

# Gradient Descent

Research Skills: Machine Learning

Grzegorz Chrupała  
g.chrupala @ uvt.nl

# Two aspects of a learner

- How it uses inputs to predict outputs
  - DT – **traverse tree** asking questions until arrive at leaf with **target**
  - Perceptron – predict **+1** if  $\mathbf{w} \cdot \mathbf{x} + b \leq 0$
- How it learns
  - DT – split data using best question and recurse on splits
  - Perceptron – update **w** and **b** if made a mistake

# Two aspects of a learner

- How it uses inputs to predict outputs
  - **Model**
- How it learns
  - **Optimization algorithm**

# Modularity

- Often **parameters** (e.g. weights) of the same model can be found in many different ways
- Standard **optimization** algorithms for many types of problems
  - Often can be treated as a **black box**

# Linear models

- Linear models are based on a weighted sum of features
- (Multiclass) Classification
- (Multivariate) Regression

**For example linear regression**

$$y = \mathbf{w} \cdot \mathbf{x} + b$$

# How can we find best **w**?

- Use specialized formula for linear regression

**OR**

- Convert into problem of finding minimum of function
- Standard solvers
  - Newton's method
  - (Stochastic) gradient descent
- **This approach works for many types of models**

# Sum of squared errors

- Want to find  $w, b$  for which error on training data is smallest
- We can use SSE as a measure of error


$$\text{SSE} = \sum_{i=1}^N (y_{\text{pred}}^i - y^i)^2$$



# Error as a function of $\mathbf{w}, b$

$$\text{Error}(\mathbf{w}, b) = \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}^i + b - y^i)^2$$


# Example – Iris



# Iris

- Find regression line which predicts **petal length** from **sepal length**
- For simplicity, fix  $b = 0$
- How does **Error( $w$ )** change as we vary  $w$ ?

Find  $w$  for which  $\text{Error}(w)$  is lowest




# Start with something even simpler

$$\text{Error}(\mathbf{w}, b) = \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}^i + b - y^i)^2$$

- Work through example of a simpler function:

$$f(w) = w^2$$




- How can we find the value of  $w$  which minimizes  $f(w)$ ?
- Start at a random value of  $w$
- Check slope of function at this point
- Descend the slope: adjust  $w$  to decrease  $f(w)$

# Gradient vs slope

- **Slope** describes steepness of a single dimension
- We usually work with functions with vectors as arguments , e.g.  $\text{Error}(\mathbf{w})$
- **Gradient** is the collection of slopes, one for each dimension





# Gradient descent for

$$f(w) = w^2$$

# How do we compute the slope of a function?

- First derivative!
- For function  $f$ , first derivative can be written  $f'$
- Then  $f'(a)$  is the slope of function  $f$  at point  $a$

# First derivative

- If we define

$$f(w) = w^2$$

- The first derivative is


$$f'(w) = 2w$$


# Ready to descend


- Initialize  $w$  to some value (e.g. 1.0)
- Update:

$$w_{\text{new}} = w_{\text{old}} - \eta \times f'(w_{\text{old}})$$


- $\eta$  is the **learning rate**, controlling speed of descent (e.g. 0.01 or 0.2)
- Stop when  $w$  doesn't change much any more

















# Back to function Error(w,b)

$$b_{\text{new}} = b_{\text{old}} - \eta \times 2 \sum_{i=1}^N (y_{\text{pred}}^i - y^i)$$



$$\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} - \eta \times 2 \sum_{i=1}^N (y_{\text{pred}}^i - y^i) \mathbf{x}^i$$



No need to memorize these  
formulas.

They are given only for  
reference.



# Visualization



(Keep fixed  $b=0$  to make it easier to plot)



























# Gradient descent

- For many types of models, we can find good model parameters with gradient descent
- Important to use appropriate learning rate!

# Problem: learning rate

- What will happen if the learning rate is too small?
- And if it's too big?



# Gradient descent with big datasets

$$\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} - \eta \times 2 \sum_{i=1}^N (y_{\text{pred}}^i - y^i) \mathbf{x}^i$$

We have to compute predictions and calculate residuals for all the examples before making an update.

# Idea: update more often

- Instead, we could update after every example (or after every 100):

1: for  $i = 1$  to  $N$  do

2:      $\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} - \eta \times 2(y_{\text{pred}}^i - y^i)\mathbf{x}^i$

# Stochastic Gradient Descent

- **Stochastic** gradient descent (aka **SGD**)
- Suitable for online learning scenarios
  - Compare with the Perceptron update rule
- Workhorse of modern Machine Learning
  - Large, deep neural networks

# Summary

- Modular learning:  
Model + Optimization
- Gradient descent to find model parameters with lowest error
- SGD for working with big data