

Logistic Regression

Research Skills: Machine Learning

Grzegorz Chrupała
[g.chrupala @ uvt.nl](mailto:g.chrupala@uvt.nl)

Models and learning algorithms

- We saw how to decouple model from learning algorithm
- (Stochastic) Gradient Descent can train various models
- Today, a classification model
 - can be fit using (S)GD

Error function

- Learning a model – minimizing error function
- Different models – different error functions
- For example, SSE for linear regression

$$\text{SSE} = \sum_{i=1}^N (y_{\text{pred}}^i - y^i)^2$$

Loss function

- Loss function quantifies our mistake on a **single example**
- **Squared loss** corresponds to **SSE**.

$$\ell_{\text{squared}}(z) = (z - y)^2$$


where $z = \mathbf{w} \cdot \mathbf{x} + b$ is the score of the model and y the target

Loss for classification

- Zero-one loss:

1 if we're made a mistake, 0 otherwise


$$\ell_{0/1}(z) = \begin{cases} 1 & \text{if } t = 1 \text{ and } z < 0 \\ 1 & \text{if } t = -1 \text{ and } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$




Loss for classification

- Zero-one loss – no gradient
- Could we just use squared loss for classification?
 - What happens if $z = 2.0$?
 - and $z = 10.0$?
 - Penalizes confident correct predictions

Example



Example



Problem

- Bad decision boundary
- Model cares too much about predicting exactly 1 for examples with high *#good*
- Need better loss function

Let's find a better way.


Regression for classifying

- In regression we predict numbers
- In classification we predict labels
- Logistic regression
 - regress on **probabilities** of labels

Logistic regression

- Let p = probability that label is positive
 - number between 0 and 1
- Logit function maps p to $[-\infty, \infty]$

$$\text{logit}(p) = \log \left(\frac{p}{1 - p} \right)$$



Examples

$$\text{logit}(0.01) = -4.6$$

$$\text{logit}(0.50) = 0.0$$

$$\text{logit}(0.99) = 4.6$$

Logistic regression


- A logistic regression model predicts $\text{logit}(p)$ using a linear model

$$\text{logit}(p)_{\text{pred}} = \mathbf{w} \cdot \mathbf{x} + b$$

Logistic regression

- We can map the logit back to probability using **inverse logit** (or **logistic**, or **sigmoid**) function

$$\text{logit}^{-1}(z) = \frac{1}{1 + \exp(-z)}$$



Examples

$$\text{logit}^{-1}(0) = 0.50$$

$$\text{logit}^{-1}(-4.6) = 0.01$$

$$\text{logit}^{-1}(4.6) = 0.99$$

Logistic regression

- Putting the pieces together

$$p_{\text{pred}} = \text{logit}^{-1}(\mathbf{w} \cdot \mathbf{x} + b)$$

Example: movie reviews

- | | #good | #dark | #mediocre | #the |
|--------------------|----------|---------------------------------|-----------|-------|
| ▪ \mathbf{x}^1 | = (1, | 0, | 0, | 5) |
| ▪ \mathbf{x}^2 | = (2, | 3, | 2, | 7) |
| ▪ \mathbf{w} | = (2.5, | 0.5, | -4.0, | 0.0) |
| ▪ b | = 0.5 | | | |
| ▪ score^1 | = 3.0, | $p^1 = \text{logit}^{-1}(3.0)$ | = 0.95 | |
| ▪ score^2 | = -1.0, | $p^2 = \text{logit}^{-1}(-1.0)$ | = 0.27 | |

Log loss function aka cross-entropy

- Loss function quantifying mistakes for LR

$$\ell_{\log}(z) = \begin{cases} -\log(p_{\text{pred}}) & \text{if } y = 1 \\ -\log(1 - p_{\text{pred}}) & \text{if } y = 0 \end{cases}$$

where $p_{\text{pred}} = \text{logit}^{-1}(z)$

- Minimize log loss – find model which gives **maximum probability** to training targets

Log loss function

$$\ell_{\log}(z) = \begin{cases} -\log(p_{\text{pred}}) & \text{if } y = 1 \\ -\log(1 - p_{\text{pred}}) & \text{if } y = 0 \end{cases}$$

where $p_{\text{pred}} = \text{logit}^{-1}(z)$

alternative notation

$$\ell_{\log}(z) = -y \log(p_{\text{pred}}) - (1 - y) \log(1 - p_{\text{pred}})$$

Summary

Logistic vs Linear Regression

Logistic vs Linear: prediction

- Both use the score of the linear model

$$z = \mathbf{w} \cdot \mathbf{x} + b$$

- Linear regression uses it directly

$$y_{\text{pred}} = z$$

- Logistic regression via inverse logit

$$p_{\text{pred}} = \text{logit}^{-1}(z)$$


Logistic vs Linear: loss

- For linear regression use squared loss

$$\ell_{\text{squared}}(z) = (z - y)^2$$

- For logistic regression use log loss

$$\ell_{\text{log}}(z) = -y \log(p_{\text{pred}}) - (1 - y) \log(1 - p_{\text{pred}})$$



Logistic vs Linear: SGD


- Both models can be learned via **SGD**
- Linear regression

$$\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} - \eta \times 2(y_{\text{pred}} - y)\mathbf{x}$$

- Logistic regression

$$\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} + \eta \times (y - p_{\text{pred}})\mathbf{x}$$

Example




Regularization for control of overfitting

- Control how hard the model tries to fit to training data
- Models with small weights less flexible – less freedom to fit data
- Penalize large weights to control overfitting

L2 regularization penalty

$$\text{Error}(\mathbf{w}, b) = \sum_{i=1}^N \ell(z^i) + \alpha \sum_{j=1}^M w_j^2$$

- We add a term to the error function which penalizes large weight
- Parameter α controls strength of penalty



Logistic regression in scikit-learn

- `sklearn.linear_model.SGDClassifier`
 - Supports several loss functions including log loss
 - Good choice for large datasets
 - Regularization via parameter **alpha**
- `sklearn.linear_model.LogisticRegression`
 - Specialized LR algorithm
 - Good for small and medium data sizes
 - Regularization via param **C=1/alpha**

Summary

- Different loss functions give rise to different linear models
- Logistic regression for probabilistic classification
- Control overfitting via L2 regularization